

Investigation into Stock Market Forecasting using Recurrent Neural Networks

Student: Oliver Crow, Supervisor: John Cartlidge, Project Type: Research

University of Bristol, Department of Computer Science

Overview

- ▶ Can RNNs (with LSTM) be used to generate accurate predictions of the stock market?
- ▶ Are the predictions reliable?
- ▶ How is accuracy and reliability of predictions affected when predicting more than 1 day into the future?
- ▶ What RNN configurations are optimal?
- ▶ How much data is required?
- ▶ What type of data best educates predictions (news, fundamentals, technical indicators, etc...)?

Project Outline

- ▶ Attempt to replicate findings and results from published articles
- ▶ Further improve findings by adding in additional items of data (e.g. news, fundamentals, technical indicators, etc...)
- ▶ Test model over configuration space
 - ▶ Number of neurons
 - ▶ Batch size
 - ▶ Epochs
 - ▶ Historical window size
 - ▶ Prediction days
- ▶ Reflect on results
- ▶ Investigate potential usefulness of predictions to assess predictive powers of the model

More Data

Giving the model a larger quantity and more diverse range of data can assist in making better predictions. Unfortunately, financial data (esp. historical) is very difficult and/or expensive. This is a large limitation in projects such as this one. The types and quality of the data fed to the network can make or break the results.

Pre-processing is a step which needs to be performed to clean the data, and make it ready for the network (e.g. fill in missing values). Textual data needs to be pre-processed heavily before being given to the network. This can include performing:

- ▶ Sentiment Analysis
- ▶ Word Embedding

Preliminary Results

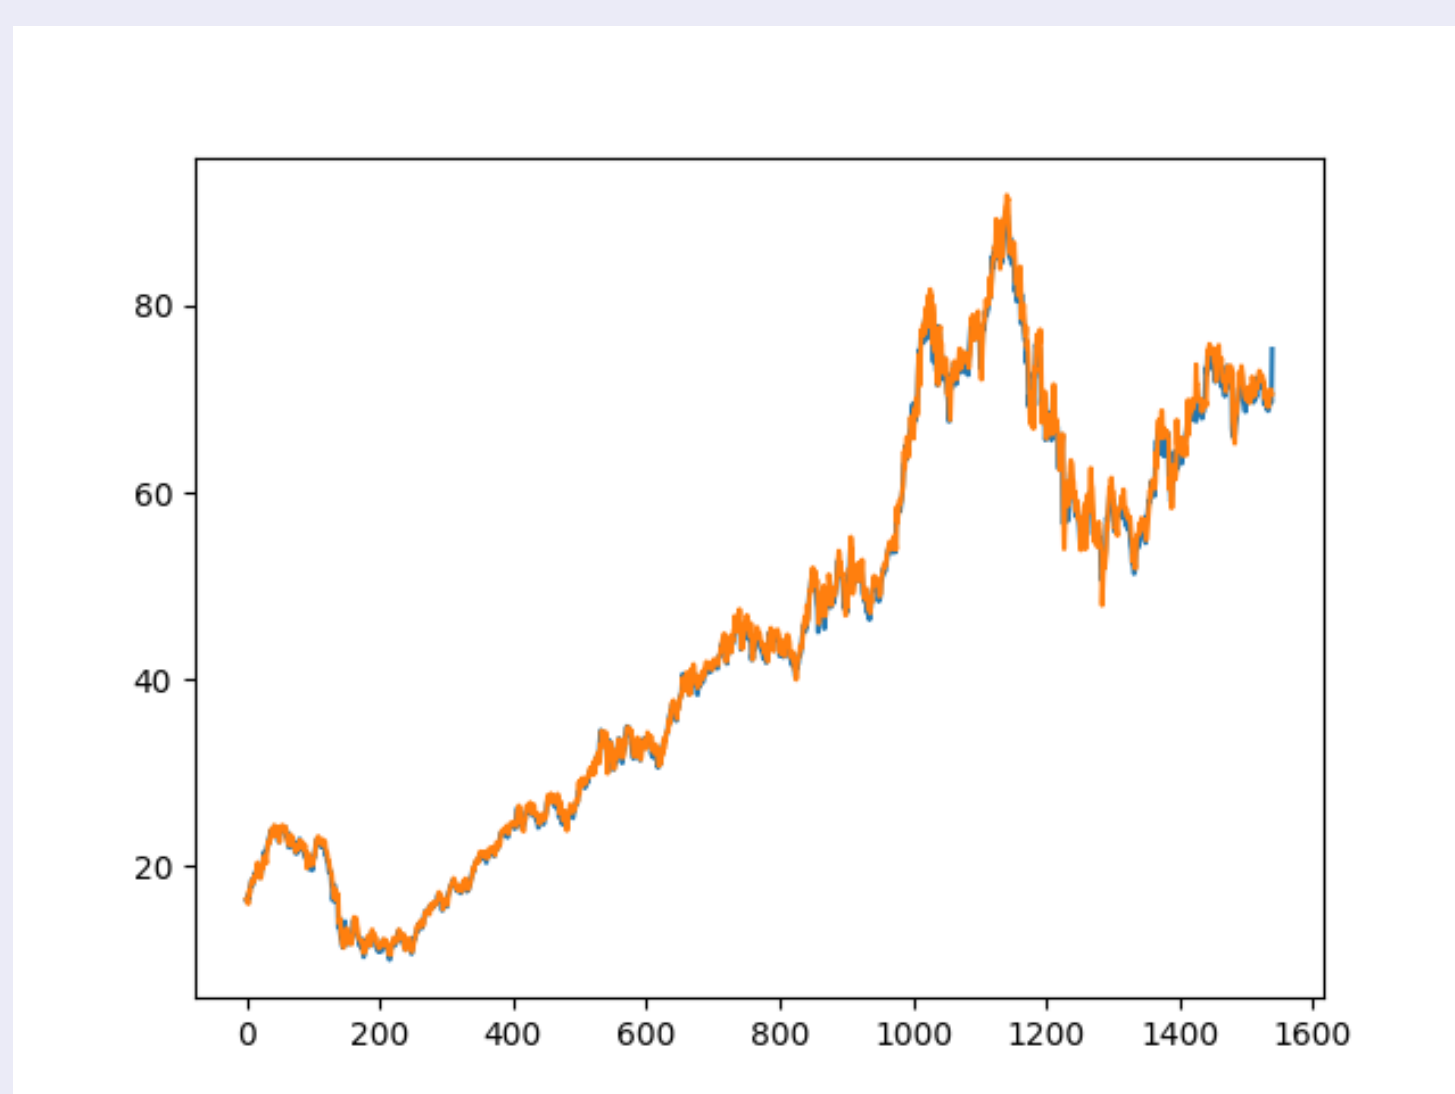


Figure: Closing Price Prediction (1-day ahead, 300-day window)
[prediction: orange, ground truth: blue]

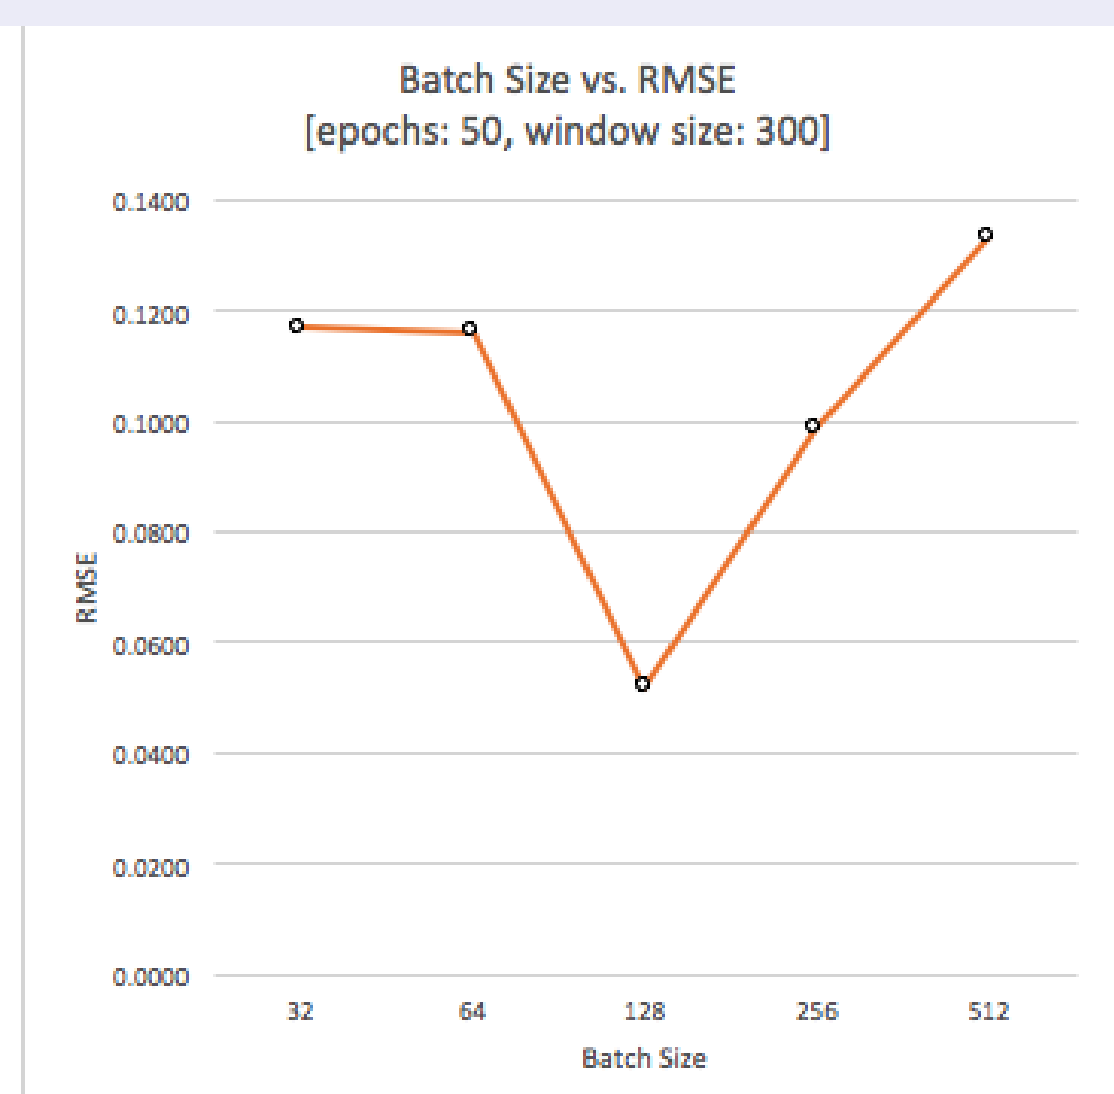



Figure: Comparison of Configurations

