


ARMv8-A Scalable Vector Extension (SVE) Simulation and Evaluation

Hal Jones, supervised by Simon McIntosh-Smith
University of Bristol, Department of Computer Science

Introduction

Conventional (**scalar**) processor instructions perform a single calculation each. Multiple calculations require multiple instructions, and thus cycles.

Vector instructions allow a processor to repeat a mathematical operation across a set of data, in a single cycle. This technique is also known as “Single Instruction, Multiple Data” (**SIMD**).

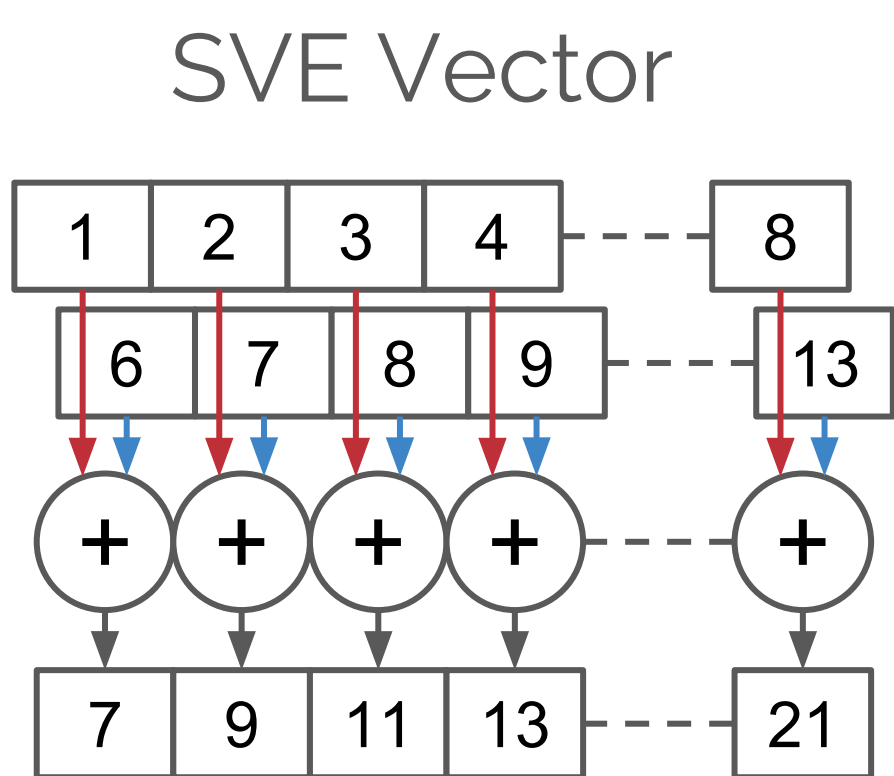


Scalable Vector Extension (SVE)

Current SIMD implementations have fixed vector lengths. Increasing the vector length requires a significant number of new instructions.

SVE is an extension to the ARMv8-A architecture that allows for hardware designers to specify the size of their vector registers. Instructions are **Vector-Length Agnostic**, and perform an operation over any vector length.

There are currently no production processors implementing SVE, so evaluating the performance of this extension is not easily possible.



Research Goals

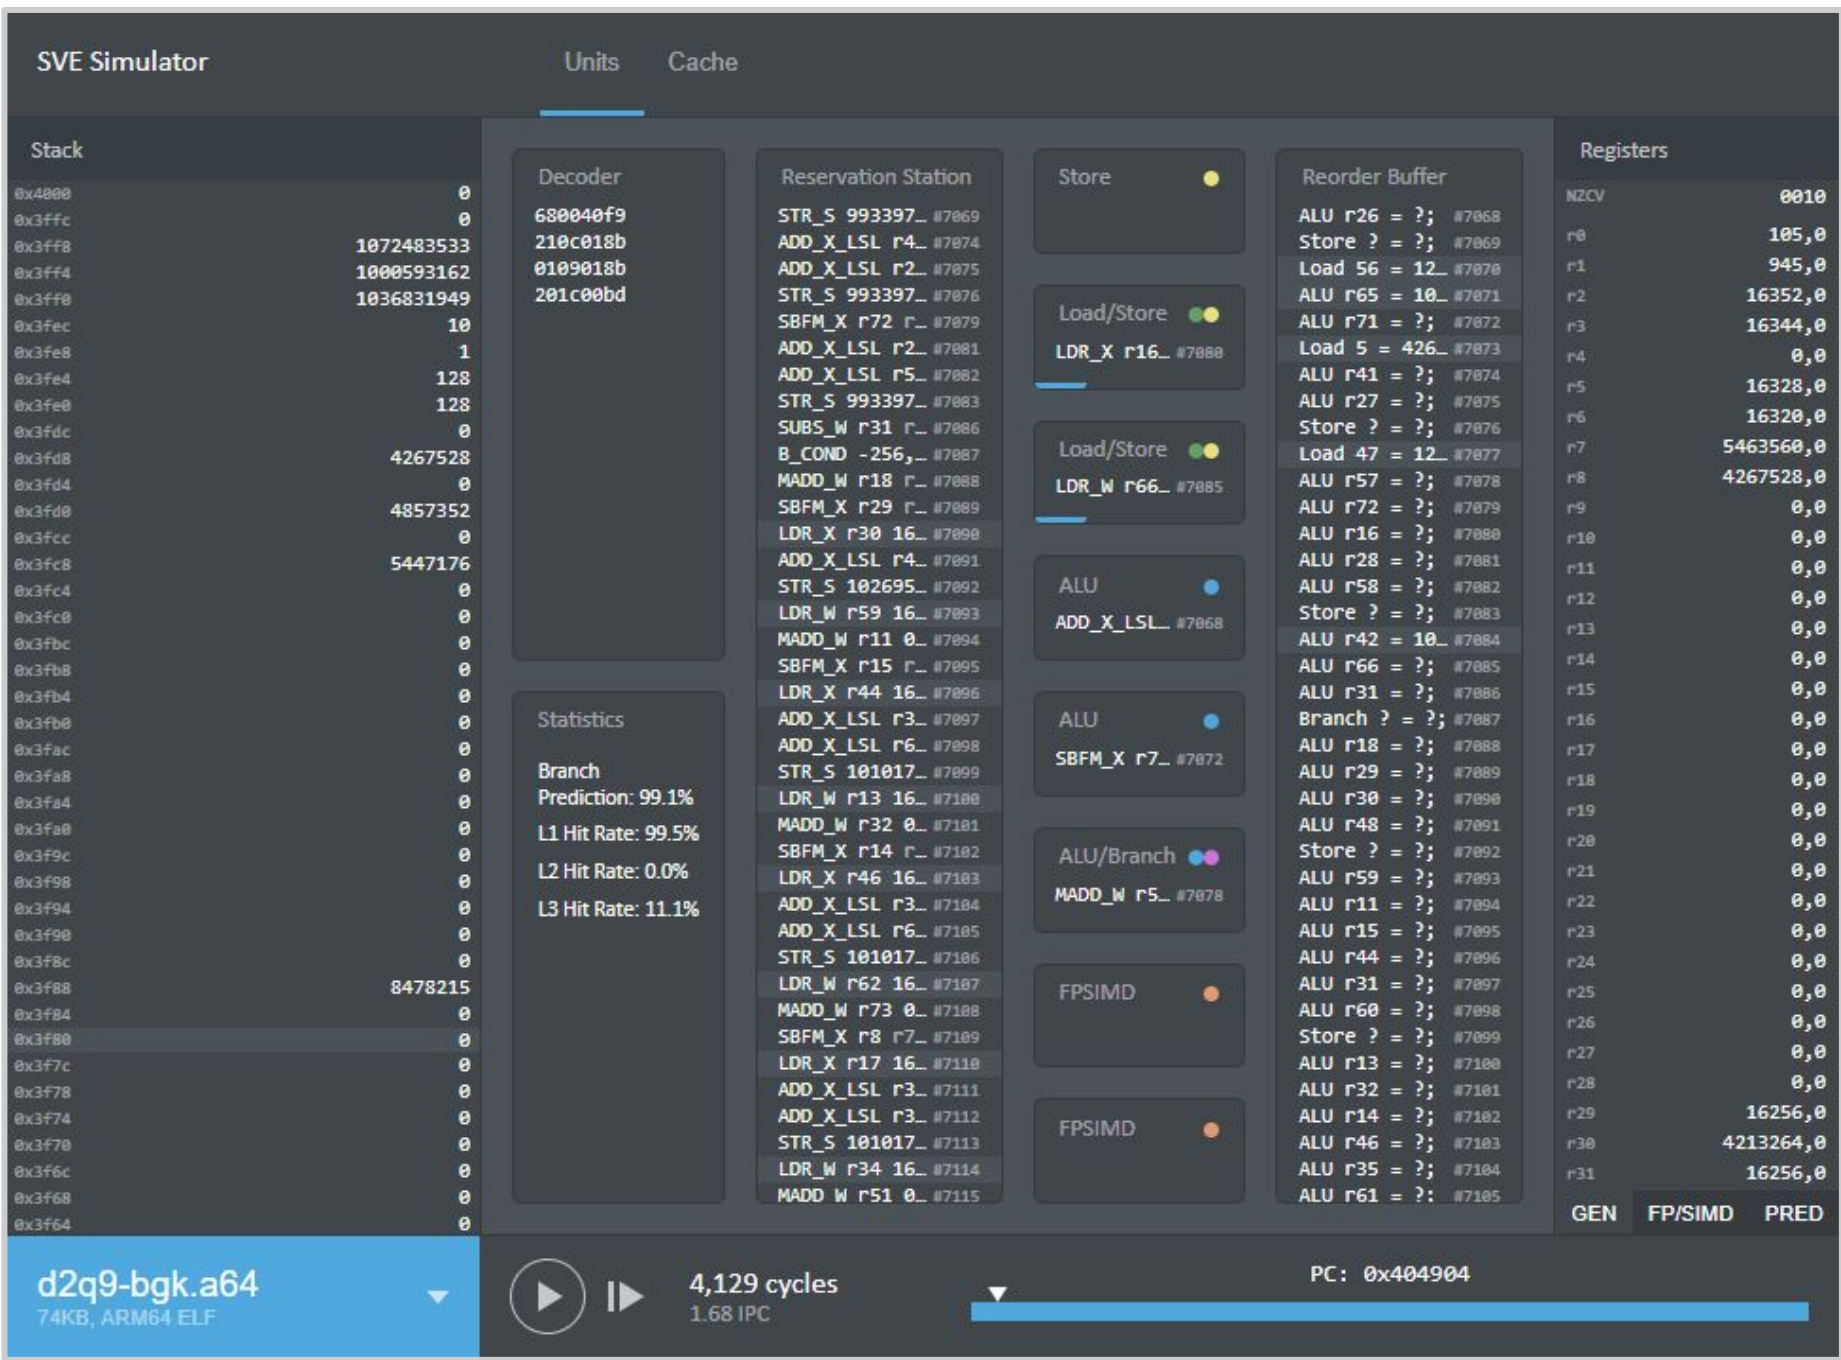
- Simulate SVE**
 - Create a **64-bit ARM processor simulator** capable of running compiled AArch64 binaries, with SVE support
- Benchmark Support**
 - Progressive implementation of simulator features will make it possible to execute **high-performance benchmarks** such as BUDE, TeaLeaf and CloverLeaf, compiled to include SVE instructions.

Performance Evaluation

Configuring the simulator to match existing an high-performance ARM processor and **measuring benchmark performance** will allow for an approximation of SVE benefits. Varying simulator properties such as cache sizes and vector length can be used to explore the benefits and limitations of SVE across a range of hardware.

Project Progress


- Processor Simulator**
 - Superscalar out-of-order with speculative execution
 - Capable of **decoding and executing AArch64 binaries**
 - Configurable multi-level cache
 - Advanced SIMD (**ASIMD**) and SVE support
 - Graphical processor/cache visualiser for debugging and execution inspection
 - Configured to match a **Cavium ThunderX2** core
- Supported Benchmarks**
 - Supports several benchmarks such as BUDE and Lattice Boltzmann, at various optimisation levels
 - Designed to be rapidly extensible to support further benchmarks



Preliminary Results

With the default 128-bit vector length, SVE performs marginally worse than ASIMD in most test benchmarks, likely due to vector predication overhead.

With a vector (and cache line) size increase, SVE code performs **significantly better** than ASIMD, with a 60% performance difference at 256 bits in the test case below. Further increasing the vector size sees diminishing returns.



Above: 1000 iterations of the *Daxpy* algorithm, compiled using *armclang*. Fewer cycles is better.